


Stage 1: controlled pretraining



random initialization;
paired L1 supervision

Stage 2: one epoch, 2,600-update cap

